

# Why is my code slow?

Salaar Liaqat

Data Sciences Institute, UofT

# Outline

- Caching, Memoization, and Vectorization
- Parallel Computing
- Greedy and Exhaustive Algorithms
- Faster Implementations versus Faster Algorithms

# Section 1

## Caching, Memoization, and Vectorization

# Caching


- *Caching* refers to storing things for later use
  - ▶ Your browser probably does by temporarily downloading page details on your local disk
  - ▶ Faster, reduces server load
  - ▶ Other examples include 3D rendering and saving common database queries
- However, caching usually takes space in exchange for faster run times
- The *space-time* trade off is a case where an algorithm trades increased space usage for faster runtimes

# Memoization

- *Memoization* refers to storing results of function calls to use for later
  - ▶ Specific method of caching
- This is useful for methods with a lot of repeated computations
- For instance, in our recursive Fibonacci number function.
- `fib(12)` is called by `fib(13)`, `fib(14)` etc.
  - ▶ And `fib(3)` is called many many times
- $F(5) = F(4) + F(3) = F(3) + F(2) + F(2) + F(1)$  Which calculates repeated subproblems

## How Memoization Works

- Since we store the results, each function call is only made once, making the time complexity  $O(n)$ , much better than  $O(2^n)$ <sup>1</sup>
- Memoization can also avoid the maximum recursion depth error because the call stack is smaller



# Memoization Python

```
cache = {0: 0, 1: 1}

def fib(n):
    if n in cache:
        return cache[n]
    else:
        cache[n] = fib(n - 1) + fib(n - 2)
        return cache[n]
```

For the base cases, we replace calling `fib(0)` and `fib(1)` by getting the values from the dictionary

# Memoization Python

- We can use the `functools` library, which is included in the standard library (no pip install needed!)
  - ▶ `functools` does memoization for you!
- We can use the `@cache` decorator, but the cached dictionary can grow to massive sizes
- Instead, `@lru_cache(maxsize = n)` uses the LRU (least recently used) `n` computations
- Alternatively, we can use `joblib` to store the memoized results in a file



# Memoization Python

```
from functools import lru_cache

@lru_cache(maxsize=10)
def fib_rec(n):
    if n == 0 or n == 1:
        return n
    else:
        return fib_rec(n-1) + fib_rec(n-2)
```

# Vectorized Operations

- *Vectorization* is a technique of implementing array operations without for loops
- We use functions defined by various modules that are highly optimized for the specific problem
- NumPy provides a lot of functions that vectorized and are faster than for loops
  - ▶ Array add/subtract/multiply/divide by scalar
  - ▶ Sum of array
  - ▶ Max/min of array
- Keep this in mind for some ML processes that are iterative, such as gradient descent

# Why Vectorized Operations Work

- Python (and R) are interpreted languages. There is no compiler and the languages are dynamic
- C language, for instance, makes optimization at the compiler level (before execution) to speed up your code
- Thus, NumPy implements arrays in C, which speeds things up
- The other reason vectorization works is because of parallelization

## Section 2

### Parallel Computing

# Parallelization

Compare the following codes. What are their run times?

```
def fib(n):  
    if n <= 1:  
        return n  
    else:  
        return fib(n - 1) + fib(n - 2)
```

# Parallelization

```
import numpy

def add_one(n, x):
    y = np.zeros(n)
    for i in range(n):
        y[i] = x[i] + 1

    return y
```

# Parallelization

- Both are  $O(n)$ , but the second code chunk can be done in *parallel* because the  $n$  computations are independent.
- Fibonacci depends on the previous two values
- The requirements for code to be parallelized and vectorized are similar, but not the same
- The Numba library can help with parallelizing your code
- Note parallel means the process takes place on one machine, but *distributed* means the computation is shared across many machines

## Section 3

### Greedy and Exhaustive Algorithms



# Greedy Approach (literally)

- Let's revisit the knapsack problem, taking a different approach.
- The items are:
  - ▶ Stereo: \$3000, 4 kg
  - ▶ Laptop: \$2000, 3 kg
  - ▶ Guitar: \$1500, 1 kg
- If we follow the rule “get the most valuable item, then get second most valuable etc.” we would make \$3000 by taking the stereo, which isn't the optimal \$3500
- A *greedy algorithm* picks the optimal move at each step, which hopefully leads to the overall optimal solution
  - ▶ But it finds the solution in  $O(n)$  time

# Greedy Approach

- Let's say you could take fractions of an item and we tried the greedy approach
  - ▶ Peanuts: \$7/kg
  - ▶ Rice: \$5/kg
  - ▶ Tea: \$12/kg
- We would take tea until it runs out, followed by peanuts and rice. This is the optimal solution in  $O(n)$  time!

# Classroom Scheduling Problem

- Suppose we want to hold as many classes in a classroom as possible <sup>2</sup>

Class	Start	End
Yoga	9AM	10AM
Music Theory	9:30AM	11AM
Painting	10AM	11AM
Algorithms	10:30AM	11:30AM
Calculus	11AM	12PM

2 Minutes: write down a greedy algorithm to solve this problem

---

<sup>2</sup>From Bhargava chapter 8

# Classroom Scheduling Problem

## Algorithm

- 1 Pick the class that ends the soonest. This is the first class you'll hold in this classroom
- 2 Now, you have to pick a class that starts after the first class. Again, pick the class that ends the soonest. This is the second class you'll hold
- 3 Repeat the second step

This not only produces the correct solution but also does so in  $O(n)$  time, for  $n$  classes!

# Classroom Scheduling Problem

- An alternative algorithm is the *exhaustive approach*
  - ▶ We try every combination of classes. At the end, we see which solution fits the most classes
  - ▶ We try every combination of items to steal. At the end, we see which solution has the most value
- While brute forcing might sound always unnecessary, there are cases where it is needed to get the optimal solution
  - ▶ When performing subset selection for regression or decision tree, we can't guarantee the variables are uncorrelated. So forward/backward stepwise selection isn't guaranteed to produce the best outcome
  - ▶ More on this in a few slides
- 2 minutes: what is the time complexity of best subset selection?

# Greedy Approximation Algorithms

- Problems involving finding the best subset of a variable to max/min an objective value are generalized as the problem of finding the best *power set*.
  - ▶ There are  $2^n$  power sets, which becomes impossible to calculate past  $n = 100$  (depending on the constants)
- *Approximation algorithms* are judged by how fast they are and how close they are to the optimal solution
  - ▶ Forward/backwards stepwise selection is an approximation algorithm to best subset selection

# N-P Complete Problems

- In the power set problem, we need to brute force all combinations and test them. Such problems are called *N-P Complete*
  - ▶ A lot of smart people think it's not possible to solve these with efficient algorithms
- It's hard to tell if a problem is N-P complete
  - ▶ Finding the shortest path between two points is N-P complete (travelling salesman)
  - ▶ But the knapsack problem isn't N-P complete because we can solve it using dynamic programming

## Section 4

### Faster Implementations versus Faster Algorithms



# Faster Implementations versus Faster Algorithms

- There are two ways we speed up our code
  - ▶ Use a faster algorithm, such as dynamic programming instead of brute force. Algorithms are concerned with the approach to the problem
  - ▶ Use a faster implementation, such as vectorization instead of loops
- It is useful to think about these separately when developing a programming, then combining them to create a super-fast approach!

## Section 5

### Recommended Problems and References

# Recommended Problems and Readings

- Cormen: Chapter 34 on NP-Completeness (highly optional)
- Bhargava: Chapter 8 exercises
  - ▶ 8.1 - 8.8
- Vectorize the second code chunk in the Parallelization section
- Find the longest palindrome from a string Hint: use a greedy algorithm
- Computing Pascal's triangle Hint: use dynamic programming

# References

- Bhargava, A. Y. (2016). *Grokking algorithms: An illustrated guide for programmers and other curious people*. Manning. Chapter 1.
- Cormen, T. H. (Ed.). (2009). *Introduction to algorithms* (3rd ed). MIT Press. Chapter 1 and 3.
-